

```
In [36]: import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.datasets import cifar10
from sklearn.model_selection import train_test_split
import tensorflow as tf
```

```
In [2]: #load the dataset
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>  
**170498071/170498071** **11s** 0us/step

```
In [3]: #Combine the train and test sets for a new 80-20 split
x_combined = np.concatenate((x_train, x_test), axis=0)
y_combined = np.concatenate((y_train, y_test), axis=0)
```

```
In [10]: #Display 5 sample images with labels
labels = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
plt.figure(figsize=(10, 2))
for i in range(5):
    plt.subplot(1, 5, i+1)
    plt.imshow(x_combined[i])
    plt.title(labels[int(y_combined[i])])
    plt.axis('off')
plt.tight_layout()
plt.show()
```

/tmp/ipython-input-10-3047913360.py:7: DeprecationWarning: Conversion of an array with ndim > 0 to a scalar is deprecated, and will error in future. Ensure you extract a single element from your array before performing this operation. (Deprecated NumPy 1.25.)

```
plt.title(labels[int(y_combined[i])])
```



```
In [11]: #Print the shape of the dataset and count unique labels
print("Shape of the dataset:", x_combined.shape)
print("Shape of the labels:", y_combined.shape)
print("Number of unique labels:", len(np.unique(y_combined)))
print("Unique labels:", np.unique(y_combined))
```

Shape of the dataset: (60000, 32, 32, 3)  
 Shape of the labels: (60000, 1)  
 Number of unique labels: 10  
 Unique labels: [0 1 2 3 4 5 6 7 8 9]

```
In [12]: #Normalise the pixel value to [0,1]
x_combined = x_combined.astype('float32') / 255.0
```

```
In [13]: #Split the train and test to (80%,20%)
x_train, x_test, y_train, y_test = train_test_split(x_combined, y_combined, test_size=0.2, random_state=42)
```

```
In [14]: print("Shape of the train set:", x_train.shape)
print("Shape of the test set:", x_test.shape)
```

Shape of the train set: (48000, 32, 32, 3)  
 Shape of the test set: (12000, 32, 32, 3)

```
In [16]: #CNN model architecture
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from tensorflow.keras.utils import to_categorical
```

```
In [17]: #convert labels to one-hot encoded vectors
y_train = to_categorical(y_train, num_classes=10)
y_test = to_categorical(y_test, num_classes=10)
```

```
In [21]: #build a simple CNN MODEL
model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)),
    MaxPooling2D((2, 2)),
    Dropout(0.25),
    Conv2D(64, (3, 3), activation='relu'),
```

```

        MaxPooling2D((2, 2)),
        Dropout(0.25),
        Flatten(),
        Dense(128, activation='relu'),
        Dropout(0.5),
        Dense(10, activation='softmax')
    ])

#Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
#Display model summary
model.summary()

```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
conv2d_4 (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d_4 (MaxPooling2D)	(None, 15, 15, 32)	0
dropout_6 (Dropout)	(None, 15, 15, 32)	0
conv2d_5 (Conv2D)	(None, 13, 13, 64)	18,496
max_pooling2d_5 (MaxPooling2D)	(None, 6, 6, 64)	0
dropout_7 (Dropout)	(None, 6, 6, 64)	0
flatten_2 (Flatten)	(None, 2304)	0
dense_4 (Dense)	(None, 128)	295,040
dropout_8 (Dropout)	(None, 128)	0
dense_5 (Dense)	(None, 10)	1,290

**Total params:** 315,722 (1.20 MB)

**Trainable params:** 315,722 (1.20 MB)

**Non-trainable params:** 0 (0.00 B)

```

In [20]: #Insights
#1.Well-structured Cnn: Conv->Pool->Dropout repeated twice, followed by Dense layers
#2.parameter-Efficient:315k parameters- suitable for CIFAR-10(smallimages)
#3.Output layer has 10 units with softmax(correct for CIFAR-10's 10 classes)

```

```

In [22]: #Train the model
history = model.fit(x_train, y_train, epochs=10, batch_size=64, validation_data=(x_test, y_test))

```

```

Epoch 1/10
750/750 ————— 66s 85ms/step - accuracy: 0.2846 - loss: 1.9258 - val_accuracy: 0.5178 - val_loss: 1.3734
Epoch 2/10
750/750 ————— 81s 84ms/step - accuracy: 0.4853 - loss: 1.4360 - val_accuracy: 0.5691 - val_loss: 1.2206
Epoch 3/10
750/750 ————— 64s 86ms/step - accuracy: 0.5361 - loss: 1.2991 - val_accuracy: 0.6046 - val_loss: 1.1514
Epoch 4/10
750/750 ————— 81s 85ms/step - accuracy: 0.5760 - loss: 1.2051 - val_accuracy: 0.6256 - val_loss: 1.0790
Epoch 5/10
750/750 ————— 64s 85ms/step - accuracy: 0.5937 - loss: 1.1549 - val_accuracy: 0.6401 - val_loss: 1.0266
Epoch 6/10
750/750 ————— 83s 86ms/step - accuracy: 0.6074 - loss: 1.1164 - val_accuracy: 0.6652 - val_loss: 0.9788
Epoch 7/10
750/750 ————— 80s 83ms/step - accuracy: 0.6178 - loss: 1.0825 - val_accuracy: 0.6684 - val_loss: 0.9339
Epoch 8/10
750/750 ————— 84s 86ms/step - accuracy: 0.6324 - loss: 1.0483 - val_accuracy: 0.6704 - val_loss: 0.9297
Epoch 9/10
750/750 ————— 80s 84ms/step - accuracy: 0.6397 - loss: 1.0221 - val_accuracy: 0.6867 - val_loss: 0.8912
Epoch 10/10
750/750 ————— 62s 83ms/step - accuracy: 0.6499 - loss: 0.9927 - val_accuracy: 0.6900 - val_loss: 0.8852

```

```


In [23]: #Insights
#we choose 10 epoch because after this layers we got a signs of overfitting so we decided to drop epoch to 10

```

```
In [24]: #Plot Accuracy and loss
plt.figure(figsize=(12, 4))
#Accuracy
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()

#Loss
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()

plt.show()
```



```
In [25]: #Evaluate the trained model
test_loss, test_acc = model.evaluate(x_test, y_test)
print('Test accuracy:', test_acc)
```

375/375 ————— 5s 13ms/step - accuracy: 0.6916 - loss: 0.8893  
Test accuracy: 0.6899999976158142

```
In [26]: #Confusion Matrix and classification report
from sklearn.metrics import confusion_matrix, classification_report
import seaborn as sns

#Get model predictions
y_pred = model.predict(x_test)
y_pred_classes = np.argmax(y_pred, axis=1)
y_test_classes = np.argmax(y_test, axis=1)


#Classification report
print(classification_report(y_test_classes, y_pred_classes))

#Confusion matrix
cm = confusion_matrix(y_test_classes, y_pred_classes)

#plot confusion matrix
plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=labels, yticklabels=labels)
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()
```

375/375 ————— 4s 10ms/step

	precision	recall	f1-score	support
0	0.75	0.71	0.73	1181
1	0.77	0.86	0.81	1188
2	0.69	0.43	0.53	1232
3	0.51	0.49	0.50	1274
4	0.62	0.64	0.63	1179
5	0.57	0.61	0.59	1230
6	0.71	0.82	0.77	1123
7	0.71	0.78	0.74	1217
8	0.83	0.79	0.81	1189
9	0.76	0.79	0.77	1187
accuracy			0.69	12000
macro avg	0.69	0.69	0.69	12000
weighted avg	0.69	0.69	0.69	12000



```
In [27]: #Insights
#1. we got an accuracy of 69
#2. from confusion matrix we can see out predictions
```

```
In [29]: #Plot
#Identify correct and incorrect predictions
correct_preds = np.where(y_pred_classes == y_test_classes)[0]
incorrect_preds = np.where(y_pred_classes != y_test_classes)[0]

#Plot 5 correctly classified images
plt.figure(figsize=(15, 5))
for i in range(5):
    plt.subplot(1, 5, i+1)
    plt.imshow(x_test[correct_preds[i]])
    plt.title(f'True: {labels[y_test_classes[correct_preds[i]]], Pred: {labels[y_pred_classes[correct_preds[i]]]}')
    plt.axis('off')

plt.tight_layout()
plt.show()
```



```
In [30]: #Top 5 incorrectly classified images
plt.figure(figsize=(15, 5))
for i in range(5):
    plt.subplot(1, 5, i+1)
    plt.imshow(x_test[incorrect_preds[i]])
    plt.title(f'True: {labels[y_test_classes[incorrect_preds[i]]], Pred: {labels[y_pred_classes[incorrect_pred:
    plt.axis('off')

plt.tight_layout()
plt.show()
```



```
In [34]: #To improve model
def build_and_train_model(optimizer, epochs=12, batch_size=64):
    model = Sequential([
        Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)),
        MaxPooling2D((2, 2)),
        Dropout(0.25),

        Conv2D(64, (3, 3), activation='relu'),
        MaxPooling2D((2, 2)),
        Dropout(0.25),

        Flatten(),
        Dense(128, activation='relu'),
        Dropout(0.5),
        Dense(10, activation='softmax')
    ])

    model.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=['accuracy'])
    history = model.fit(x_train, y_train, epochs=epochs, batch_size=batch_size, validation_data=(x_test, y_test)

    test_loss, test_acc = model.evaluate(x_test, y_test)
    print(f'Test accuracy with {optimizer}: {test_acc}')

    return history, test_acc
```

```
In [39]: #run with different optimizers
optimizers={
    'Adam':tf.keras.optimizers.Adam(),
    'SGD':tf.keras.optimizers.SGD(),
    'RMSprop':tf.keras.optimizers.RMSprop()
}
results={}
for name,opt in optimizers.items():
    history,acc=build_and_train_model(opt)
    results[name]=acc
```

```
Epoch 1/12
750/750 ————— 68s 89ms/step - accuracy: 0.3008 - loss: 1.8891 - val_accuracy: 0.5120 - val_loss: 1.3632
Epoch 2/12
750/750 ————— 80s 86ms/step - accuracy: 0.4868 - loss: 1.4217 - val_accuracy: 0.5867 - val_loss: 1.1750
Epoch 3/12
750/750 ————— 83s 88ms/step - accuracy: 0.5430 - loss: 1.2814 - val_accuracy: 0.6208 - val_loss: 1.0856
```

Epoch 4/12  
750/750 — 82s 88ms/step - accuracy: 0.5810 - loss: 1.1920 - val\_accuracy: 0.6452 - val\_loss: 1.0219

Epoch 5/12  
750/750 — 81s 87ms/step - accuracy: 0.5947 - loss: 1.1397 - val\_accuracy: 0.6566 - val\_loss: 0.9747

Epoch 6/12  
750/750 — 80s 85ms/step - accuracy: 0.6164 - loss: 1.0964 - val\_accuracy: 0.6688 - val\_loss: 0.9439

Epoch 7/12  
750/750 — 64s 85ms/step - accuracy: 0.6265 - loss: 1.0710 - val\_accuracy: 0.6758 - val\_loss: 0.9220

Epoch 8/12  
750/750 — 65s 87ms/step - accuracy: 0.6409 - loss: 1.0331 - val\_accuracy: 0.6799 - val\_loss: 0.9093

Epoch 9/12  
750/750 — 79s 83ms/step - accuracy: 0.6544 - loss: 0.9880 - val\_accuracy: 0.6942 - val\_loss: 0.8813

Epoch 10/12  
750/750 — 63s 84ms/step - accuracy: 0.6546 - loss: 0.9851 - val\_accuracy: 0.7018 - val\_loss: 0.8601

Epoch 11/12  
750/750 — 84s 87ms/step - accuracy: 0.6683 - loss: 0.9544 - val\_accuracy: 0.6852 - val\_loss: 0.8916

Epoch 12/12  
750/750 — 64s 85ms/step - accuracy: 0.6674 - loss: 0.9540 - val\_accuracy: 0.7082 - val\_loss: 0.8447

375/375 — 5s 14ms/step - accuracy: 0.7083 - loss: 0.8441

Test accuracy with <keras.src.optimizers.adam.Adam object at 0x788983838cd0>: 0.7082499861717224

Epoch 1/12  
750/750 — 65s 86ms/step - accuracy: 0.1325 - loss: 2.2787 - val\_accuracy: 0.2496 - val\_loss: 2.0588

Epoch 2/12  
750/750 — 81s 84ms/step - accuracy: 0.2437 - loss: 2.0542 - val\_accuracy: 0.3185 - val\_loss: 1.9199

Epoch 3/12  
750/750 — 84s 86ms/step - accuracy: 0.2953 - loss: 1.9238 - val\_accuracy: 0.3629 - val\_loss: 1.7846

Epoch 4/12  
750/750 — 80s 84ms/step - accuracy: 0.3390 - loss: 1.8089 - val\_accuracy: 0.4011 - val\_loss: 1.6952

Epoch 5/12  
750/750 — 82s 84ms/step - accuracy: 0.3715 - loss: 1.7284 - val\_accuracy: 0.4218 - val\_loss: 1.6093

Epoch 6/12  
750/750 — 84s 85ms/step - accuracy: 0.3955 - loss: 1.6629 - val\_accuracy: 0.4565 - val\_loss: 1.5366

Epoch 7/12  
750/750 — 66s 88ms/step - accuracy: 0.4158 - loss: 1.6039 - val\_accuracy: 0.4717 - val\_loss: 1.4806

Epoch 8/12  
750/750 — 78s 83ms/step - accuracy: 0.4317 - loss: 1.5713 - val\_accuracy: 0.4866 - val\_loss: 1.4392

Epoch 9/12  
750/750 — 83s 84ms/step - accuracy: 0.4529 - loss: 1.5206 - val\_accuracy: 0.4932 - val\_loss: 1.4275

Epoch 10/12  
750/750 — 63s 84ms/step - accuracy: 0.4631 - loss: 1.4947 - val\_accuracy: 0.5202 - val\_loss: 1.3753

Epoch 11/12  
750/750 — 64s 86ms/step - accuracy: 0.4701 - loss: 1.4726 - val\_accuracy: 0.5225 - val\_loss: 1.3625

Epoch 12/12  
750/750 — 82s 86ms/step - accuracy: 0.4816 - loss: 1.4455 - val\_accuracy: 0.5290 - val\_loss: 1.3381

375/375 — 6s 15ms/step - accuracy: 0.5253 - loss: 1.3408

Test accuracy with <keras.src.optimizers.sgd.SGD object at 0x7889be0d2650>: 0.5289999842643738

Epoch 1/12  
750/750 — 66s 86ms/step - accuracy: 0.2953 - loss: 1.9259 - val\_accuracy: 0.5207 - val\_loss: 1.3499

Epoch 2/12  
750/750 — 83s 87ms/step - accuracy: 0.4837 - loss: 1.4506 - val\_accuracy: 0.5706 - val\_loss: 1.2156

Epoch 3/12  
750/750 — 80s 85ms/step - accuracy: 0.5384 - loss: 1.3067 - val\_accuracy: 0.6172 - val\_loss: 1.1173

Epoch 4/12  
750/750 — 64s 85ms/step - accuracy: 0.5796 - loss: 1.2010 - val\_accuracy: 0.6301 - val\_loss: 1.0592

Epoch 5/12  
750/750 — 84s 87ms/step - accuracy: 0.5981 - loss: 1.1521 - val\_accuracy: 0.6533 - val\_loss: 1.0201

Epoch 6/12

750/750 — 84s 89ms/step - accuracy: 0.6141 - loss: 1.1098 - val\_accuracy: 0.6597 - val\_loss: 1.0207  
 Epoch 7/12  
 750/750 — 79s 86ms/step - accuracy: 0.6290 - loss: 1.0789 - val\_accuracy: 0.6678 - val\_loss: 0.9761  
 Epoch 8/12  
 750/750 — 82s 86ms/step - accuracy: 0.6401 - loss: 1.0496 - val\_accuracy: 0.6447 - val\_loss: 1.0299  
 Epoch 9/12  
 750/750 — 64s 85ms/step - accuracy: 0.6471 - loss: 1.0350 - val\_accuracy: 0.6720 - val\_loss: 0.9424  
 Epoch 10/12  
 750/750 — 83s 87ms/step - accuracy: 0.6483 - loss: 1.0289 - val\_accuracy: 0.6404 - val\_loss: 1.0699  
 Epoch 11/12  
 750/750 — 64s 85ms/step - accuracy: 0.6592 - loss: 1.0143 - val\_accuracy: 0.6900 - val\_loss: 0.9521  
 Epoch 12/12  
 750/750 — 82s 86ms/step - accuracy: 0.6580 - loss: 1.0011 - val\_accuracy: 0.7001 - val\_loss: 0.9029  
 375/375 — 4s 12ms/step - accuracy: 0.6983 - loss: 0.9039  
 Test accuracy with <keras.src.optimizers.rmsprop.RMSprop object at 0x788983f47890>: 0.700083315372467

```
In [40]: #performance comparison table
import pandas as pd
results_df=pd.DataFrame(list(results.items()),columns=['Optimizer','Test Accuracy'])
results_df
```

```
Out[40]:
```

	Optimizer	Test Accuracy
0	Adam	0.708250
1	SGD	0.529000
2	RMSprop	0.700083

```
In [41]: #Insights
#1.adam generally gives the best accuracy for CNNs on CIFAR-10 due to its adaptive learning rate
#2.SGD lags in performance
#3.RMSprop is also performing well
```